

```
from google.colab import drive
drive.mount('/content/gdrive', force_remount= True)
```

```
Mounted at /content/gdrive
```

```
# !pwd
%cd gdrive/My Drive/face/face
# !ls
```

```
/content/gdrive/My Drive/face/face
```

First of all, let's import necessary Packages

```
import numpy as np
import pandas as pd
import os
import matplotlib.pyplot as plt
%matplotlib inline
from math import sin, cos, pi
import cv2
from tqdm.notebook import tqdm

from keras.layers.advanced_activations import LeakyReLU
from keras.models import Sequential, Model, load_model
from keras.layers import Activation, Convolution2D, MaxPooling2D, BatchNormalization, Flatten
from keras.callbacks import ReduceLROnPlateau, ModelCheckpoint
from keras.optimizers import Adam
```

▼ Input data

```
train_data = pd.read_csv('training.csv')
test_data = pd.read_csv('test.csv')
idlookup_data = pd.read_csv('IdLookupTable.csv')

# Explore the data
train_data.head().T
```

	0	1	2	3	4
left_eye_center_x	66.0336	64.3329	65.0571	65.2257	66.7253
left_eye_center_y	39.0023	34.9701	34.9096	37.2618	39.6213
right_eye_center_x	30.227	29.9493	30.9038	32.0231	32.2448
right_eye_center_y	36.4217	33.4487	34.9096	37.2618	38.042
left_eye_inner_corner_x	59.5821	58.8562	59.412	60.0033	58.5659
left_eye_inner_corner_y	39.6474	35.2743	36.321	39.1272	39.6213
left_eye_outer_corner_x	73.1303	70.7227	70.9844	72.3147	72.5159
left_eye_outer_corner_y	39.97	36.1872	36.321	38.381	39.8845
right_eye_inner_corner_x	36.3566	36.0347	37.6781	37.6186	36.9824
right_eye_inner_corner_y	37.3894	34.3615	36.321	38.7541	39.0949
right_eye_outer_corner_x	23.4529	24.4725	24.9764	25.3073	22.5061
right_eye_outer_corner_y	37.3894	33.1444	36.6032	38.0079	38.3052
left_eyebrow_inner_end_x	56.9533	53.9874	55.7425	56.4338	57.2496
left_eyebrow_inner_end_y	29.0336	28.2759	27.5709	30.9299	30.6722
left_eyebrow_outer_end_x	80.2271	78.6342	78.8874	77.9103	77.7629
left_eyebrow_outer_end_y	32.2281	30.4059	32.6516	31.6657	31.7372
right_eyebrow_inner_end_x	40.2276	42.7289	42.1939	41.6715	38.0354
right_eyebrow_inner_end_y	29.0023	26.146	28.1355	31.05	30.9354
right_eyebrow_outer_end_x	16.3564	16.8654	16.7912	20.458	15.9259
right_eyebrow_outer_end_y	29.6475	27.0589	32.0871	29.9093	30.6722
nose_tip_x	44.4206	48.2063	47.5573	51.8851	43.2995
nose_tip_y	57.0668	55.6609	53.5389	54.1665	64.8895
mouth_left_corner_x	61.1953	56.4214	60.8229	65.5989	60.6714
mouth_left_corner_y	79.9702	76.352	73.0143	72.7037	77.5232
mouth_right_corner_x	28.6145	35.1224	33.7263	37.2455	31.1918
mouth_right_corner_y	77.389	76.0477	72.732	74.1955	76.9973
mouth_center_top_lip_x	43.3126	46.6846	47.2749	50.3032	44.9627
mouth_center_top_lip_y	70.0055	70.0000	70.4010	70.0017	70.7074

```
test_data.head()
```

	ImageId	Image
0	1	182 183 182 182 180 180 176 169 156 137 124 10...
1	2	76 87 81 72 65 59 64 76 69 42 31 38 49 58 58 4...
2	3	177 176 174 170 169 169 168 166 166 166 161 14...
3	4	176 174 174 175 174 174 176 176 175 171 165 15...

```
idlookup_data.head().T
```

	0	1	2	3
RowId	1	2	3	4
ImageId	1	1	1	1
FeatureName	left_eye_center_x	left_eye_center_y	right_eye_center_x	right_eye_center_y
Location	NaN	NaN	NaN	NaN

Check for any images with missing pixel values

```
print("Length of train data: {}".format(len(train_data)))
print("Number of Images with missing pixel values: {}".format(len(train_data) - int(train_data['Location'].isnull().sum())))

Length of train data: 7049
Number of Images with missing pixel values: 0
```

Find columns having Null values and their counts

```
train_data.isnull().sum()

left_eye_center_x      10
left_eye_center_y      10
right_eye_center_x     13
right_eye_center_y     13
left_eye_inner_corner_x 4778
left_eye_inner_corner_y 4778
left_eye_outer_corner_x 4782
left_eye_outer_corner_y 4782
right_eye_inner_corner_x 4781
right_eye_inner_corner_y 4781
right_eye_outer_corner_x 4781
right_eye_outer_corner_y 4781
left_eyebrow_inner_end_x 4779
left_eyebrow_inner_end_y 4779
left_eyebrow_outer_end_x 4824
left_eyebrow_outer_end_y 4824
right_eyebrow_inner_end_x 4779
right_eyebrow_inner_end_y 4779
```

```

right_eyebrow_outer_end_x    4813
right_eyebrow_outer_end_y    4813
nose_tip_x                    0
nose_tip_y                    0
mouth_left_corner_x          4780
mouth_left_corner_y          4780
mouth_right_corner_x          4779
mouth_right_corner_y          4779
mouth_center_top_lip_x        4774
mouth_center_top_lip_y        4774
mouth_center_bottom_lip_x     33
mouth_center_bottom_lip_y     33
Image                          0
dtype: int64

```

Define a function to plot facial keypoints with images

```

def plot_sample(image, keypoint, axis, title):
    image = image.reshape(96,96)
    axis.imshow(image, cmap='gray')
    axis.scatter(keypoint[0::2], keypoint[1::2], marker='x', s=20)
    plt.title(title)

```

Separate data into clean & unclean subsets

```
%%time
```

```

clean_train_data = train_data.dropna()
print("clean_train_data shape: {}".format(np.shape(clean_train_data)))

```

```

# https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.fillna.html
unclean_train_data = train_data.fillna(method = 'ffill')
print("unclean_train_data shape: {}\n".format(np.shape(unclean_train_data)))

```

```

clean_train_data shape: (2140, 31)
unclean_train_data shape: (7049, 31)

```

```

CPU times: user 14.1 ms, sys: 5.28 ms, total: 19.4 ms
Wall time: 21.2 ms

```

```

include_unclean_data = True    # Whether to include samples with missing keypoint values. Not
sample_image_index = 20       # Index of sample train image used for visualizing various augment

```

```
%%time
```

```

def load_images(image_data):
    images = []
    for idx, sample in image_data.iterrows():
        image = np.array(sample['Image'].split(' '), dtype=int)

```

```
        image = np.reshape(image, (96,96,1))
        images.append(image)
    images = np.array(images)/255.
    return images

def load_keypoints(keypoint_data):
    keypoint_data = keypoint_data.drop('Image',axis = 1)
    keypoint_features = []
    for idx, sample_keypoints in keypoint_data.iterrows():
        keypoint_features.append(sample_keypoints)
    keypoint_features = np.array(keypoint_features, dtype = 'float')
    return keypoint_features

clean_train_images = load_images(clean_train_data)
print("Shape of clean_train_images: {}".format(np.shape(clean_train_images)))
clean_train_keypoints = load_keypoints(clean_train_data)
print("Shape of clean_train_keypoints: {}".format(np.shape(clean_train_keypoints)))
test_images = load_images(test_data)
print("Shape of test_images: {}".format(np.shape(test_images)))

train_images = clean_train_images
train_keypoints = clean_train_keypoints
fig, axis = plt.subplots()
plot_sample(clean_train_images[sample_image_index], clean_train_keypoints[sample_image_index])

if include_unclean_data:
    unclean_train_images = load_images(unclean_train_data)
    print("Shape of unclean_train_images: {}".format(np.shape(unclean_train_images)))
    unclean_train_keypoints = load_keypoints(unclean_train_data)
    print("Shape of unclean_train_keypoints: {}\n".format(np.shape(unclean_train_keypoints)))
    train_images = np.concatenate((train_images, unclean_train_images))
    train_keypoints = np.concatenate((train_keypoints, unclean_train_keypoints))
```

```

Shape of clean_train_images: (2140, 96, 96, 1)
Shape of clean_train_keypoints: (2140, 30)
Shape of test_images: (1783, 96, 96, 1)
Shape of unclean_train_images: (7049, 96, 96, 1)

```

We can observe that approx. 68% of data is missing for several keypoints, so we need to clean the data.

```

wall time: 12.3 s

%%time

clean_train_data = train_data.dropna()
print("clean_train_data shape: {}".format(np.shape(clean_train_data)))

# https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.fillna.html
unclean_train_data = train_data.fillna(method = 'ffill')
print("unclean_train_data shape: {}\n".format(np.shape(unclean_train_data)))

clean_train_data shape: (2140, 31)
unclean_train_data shape: (7049, 31)

CPU times: user 12.6 ms, sys: 0 ns, total: 12.6 ms
Wall time: 12.3 ms

```

▼ Data Augmentation

There are various augmentation choices

```

horizontal_flip = False
rotation_augmentation = True
brightness_augmentation = True
shift_augmentation = True
random_noise_augmentation = True

rotation_angles = [12]      # Rotation angle in degrees (includes both clockwise & anti-clockwi
pixel_shifts = [12]        # Horizontal & vertical shift amount in pixels (includes shift from al

NUM_EPOCHS = 80
BATCH_SIZE = 64

```

▼ 1. Performing Horizontal Flipping for Data Augmentation

```

def left_right_flip(images, keypoints):
    flipped_keypoints = []
    flipped_images = np.flip(images, axis=2)    # Flip column-wise (axis=2)

```

```

for idx, sample_keypoints in enumerate(keypoints):
    flipped_keypoints.append([96.-coor if idx%2==0 else coor for idx,coor in enumerate(sample_keypoints)])
return flipped_images, flipped_keypoints

if horizontal_flip:
    flipped_train_images, flipped_train_keypoints = left_right_flip(clean_train_images, clean_train_keypoints)
    print("Shape of flipped_train_images: {}".format(np.shape(flipped_train_images)))
    print("Shape of flipped_train_keypoints: {}".format(np.shape(flipped_train_keypoints)))
    train_images = np.concatenate((train_images, flipped_train_images))
    train_keypoints = np.concatenate((train_keypoints, flipped_train_keypoints))
    fig, axis = plt.subplots()
    plot_sample(flipped_train_images[sample_image_index], flipped_train_keypoints[sample_image_index])

```

▼ 2. Performing Rotation Augmentation

```

def rotate_augmentation(images, keypoints):
    rotated_images = []
    rotated_keypoints = []
    print("Augmenting for angles (in degrees): ")
    for angle in rotation_angles: # Rotation augmentation for a list of angle values
        for angle in [angle, -angle]:
            print(f'{angle}', end=' ')
            M = cv2.getRotationMatrix2D((48,48), angle, 1.0)
            angle_rad = -angle*pi/180. # Obtain angle in radians from angle in degrees (n)
            # For train_images
            for image in images:
                rotated_image = cv2.warpAffine(image, M, (96,96), flags=cv2.INTER_CUBIC)
                rotated_images.append(rotated_image)
            # For train_keypoints
            for keypoint in keypoints:
                rotated_keypoint = keypoint - 48. # Subtract the middle value of the image
                for idx in range(0, len(rotated_keypoint), 2):
                    # https://in.mathworks.com/matlabcentral/answers/93554-how-can-i-rotate-a
                    rotated_keypoint[idx] = rotated_keypoint[idx]*cos(angle_rad)-rotated_keypoint[idx+1]*sin(angle_rad)
                    rotated_keypoint[idx+1] = rotated_keypoint[idx]*sin(angle_rad)+rotated_keypoint[idx+1]*cos(angle_rad)
                rotated_keypoint += 48. # Add the earlier subtracted value
                rotated_keypoints.append(rotated_keypoint)

    return np.reshape(rotated_images, (-1,96,96,1)), rotated_keypoints

if rotation_augmentation:
    rotated_train_images, rotated_train_keypoints = rotate_augmentation(clean_train_images, clean_train_keypoints)
    print("\nShape of rotated_train_images: {}".format(np.shape(rotated_train_images)))
    print("Shape of rotated_train_keypoints: {}".format(np.shape(rotated_train_keypoints)))
    train_images = np.concatenate((train_images, rotated_train_images))
    train_keypoints = np.concatenate((train_keypoints, rotated_train_keypoints))
    fig, axis = plt.subplots()
    plot_sample(rotated_train_images[sample_image_index], rotated_train_keypoints[sample_image_index])

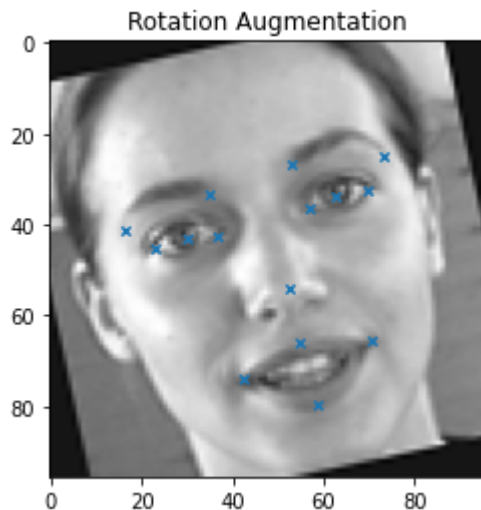
```

Augmenting for angles (in degrees):

12 -12

Shape of rotated_train_images: (4280, 96, 96, 1)

Shape of rotated_train_keypoints: (4280, 30)



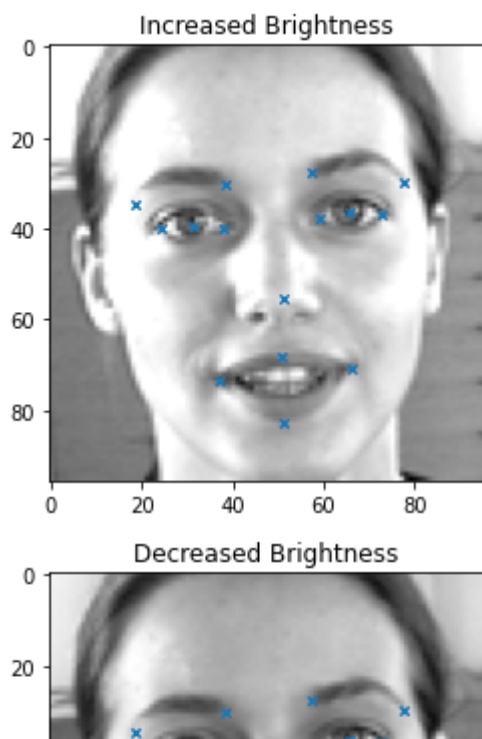
▼ 3. Performing Brightness Alteration for Data Augmentation

```
def alter_brightness(images, keypoints):
    altered_brightness_images = []
    inc_brightness_images = np.clip(images*1.2, 0.0, 1.0)    # Increased brightness by a fact
    dec_brightness_images = np.clip(images*0.6, 0.0, 1.0)    # Decreased brightness by a fact
    altered_brightness_images.extend(inc_brightness_images)
    altered_brightness_images.extend(dec_brightness_images)
    return altered_brightness_images, np.concatenate((keypoints, keypoints))

if brightness_augmentation:
    altered_brightness_train_images, altered_brightness_train_keypoints = alter_brightness(cl
    print(f"Shape of altered_brightness_train_images: {np.shape(altered_brightness_train_imag
    print(f"Shape of altered_brightness_train_keypoints: {np.shape(altered_brightness_train_k
    train_images = np.concatenate((train_images, altered_brightness_train_images))
    train_keypoints = np.concatenate((train_keypoints, altered_brightness_train_keypoints))
    fig, axis = plt.subplots()
    plot_sample(altered_brightness_train_images[sample_image_index], altered_brightness_train
    fig, axis = plt.subplots()
    plot_sample(altered_brightness_train_images[len(altered_brightness_train_images)//2+sampl
```


Shape of altered_brightness_train_images: (4280, 96, 96, 1)

Shape of altered_brightness_train_keypoints: (4280, 30)



▼ 4. Performing Horizontal & Vertical shift




```
def shift_images(images, keypoints):
    shifted_images = []
    shifted_keypoints = []
    for shift in pixel_shifts:      # Augmenting over several pixel shift values
        for (shift_x, shift_y) in [(-shift, -shift), (-shift, shift), (shift, -shift), (shift, shift)]:
            M = np.float32([[1, 0, shift_x], [0, 1, shift_y]])
            for image, keypoint in zip(images, keypoints):
                shifted_image = cv2.warpAffine(image, M, (96, 96), flags=cv2.INTER_CUBIC)
                shifted_keypoint = np.array([(point + shift_x) if idx % 2 == 0 else (point + shift_y)
                                              if np.all(0.0 < shifted_keypoint) and np.all(shifted_keypoint < 96.0):
                shifted_images.append(shifted_image.reshape(96, 96, 1))
                shifted_keypoints.append(shifted_keypoint)
    shifted_keypoints = np.clip(shifted_keypoints, 0.0, 96.0)
    return shifted_images, shifted_keypoints

if shift_augmentation:
    shifted_train_images, shifted_train_keypoints = shift_images(clean_train_images, clean_train_keypoints)
    print(f"Shape of shifted_train_images: {np.shape(shifted_train_images)}")
    print(f"Shape of shifted_train_keypoints: {np.shape(shifted_train_keypoints)}")
    train_images = np.concatenate((train_images, shifted_train_images))
    train_keypoints = np.concatenate((train_keypoints, shifted_train_keypoints))
    fig, axis = plt.subplots()
    plot_sample(shifted_train_images[sample_image_index], shifted_train_keypoints[sample_image_index])
```

Shape of shifted_train_images: (6350, 96, 96, 1)

Shape of shifted_train_keypoints: (6350, 30)




▼ 5. Adding Random Noise for Data Augmentation

```
def add_noise(images):
    noisy_images = []
    for image in images:
        noisy_image = cv2.add(image, 0.008*np.random.randn(96,96,1))    # Adding random norma
        noisy_images.append(noisy_image.reshape(96,96,1))
    return noisy_images
```

```
if random_noise_augmentation:
    noisy_train_images = add_noise(clean_train_images)
    print(f"Shape of noisy_train_images: {np.shape(noisy_train_images)}")
    train_images = np.concatenate((train_images, noisy_train_images))
    train_keypoints = np.concatenate((train_keypoints, clean_train_keypoints))
    fig, axis = plt.subplots()
    plot_sample(noisy_train_images[sample_image_index], clean_train_keypoints[sample_image_in
```

Shape of noisy_train_images: (2140, 96, 96, 1)



Visualize Train images & corresponding Keypoints

```

print("Shape of final train_images: {}".format(np.shape(train_images)))
print("Shape of final train_keypoints: {}".format(np.shape(train_keypoints)))

print("\n Clean Train Data: ")
fig = plt.figure(figsize=(20,8))
for i in range(10):
    axis = fig.add_subplot(2, 5, i+1, xticks=[], yticks=[])
    plot_sample(clean_train_images[i], clean_train_keypoints[i], axis, "")
plt.show()

if include_unclean_data:
    print("Unclean Train Data: ")
    fig = plt.figure(figsize=(20,8))
    for i in range(10):
        axis = fig.add_subplot(2, 5, i+1, xticks=[], yticks=[])
        plot_sample(unclean_train_images[i], unclean_train_keypoints[i], axis, "")
    plt.show()

if horizontal_flip:
    print("Horizontal Flip Augmentation: ")
    fig = plt.figure(figsize=(20,8))
    for i in range(10):
        axis = fig.add_subplot(2, 5, i+1, xticks=[], yticks=[])
        plot_sample(flipped_train_images[i], flipped_train_keypoints[i], axis, "")
    plt.show()

if rotation_augmentation:
    print("Rotation Augmentation: ")
    fig = plt.figure(figsize=(20,8))
    for i in range(10):
        axis = fig.add_subplot(2, 5, i+1, xticks=[], yticks=[])
        plot_sample(rotated_train_images[i], rotated_train_keypoints[i], axis, "")
    plt.show()

if brightness_augmentation:
    print("Brightness Augmentation: ")
    fig = plt.figure(figsize=(20,8))
    for i in range(10):
        axis = fig.add_subplot(2, 5, i+1, xticks=[], yticks=[])
        plot_sample(altered_brightness_train_images[i], altered_brightness_train_keypoints[i], axis, "")
    plt.show()

if shift_augmentation:
    print("Shift Augmentation: ")
    fig = plt.figure(figsize=(20,8))
    for i in range(10):
        axis = fig.add_subplot(2, 5, i+1, xticks=[], yticks=[])
        plot_sample(shifted_train_images[i], shifted_train_keypoints[i], axis, "")

```


```
plt.show()

if random_noise_augmentation:
    print("Random Noise Augmentation: ")
    fig = plt.figure(figsize=(20,8))
    for i in range(10):
        axis = fig.add_subplot(2, 5, i+1, xticks=[], yticks=[])
        plot_sample(noisy_train_images[i], clean_train_keypoints[i], axis, "")
    plt.show()
```


Shape of final_train_images: (26239, 96, 96, 1)

Shape of final_train_keypoints: (26239, 30)


Clean Train Data:



Unclean Train Data:



Rotation Augmentation:




Brightness Augmentation:






Shift Augmentation:



Random Noise Augmentation:



▼ Model

```
model = Sequential()

# Input dimensions: (None, 96, 96, 1)
model.add(Convolution2D(32, (3,3), padding='same', use_bias=False, input_shape=(96,96,1)))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())
# Input dimensions: (None, 96, 96, 32)
model.add(Convolution2D(32, (3,3), padding='same', use_bias=False))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())
model.add(MaxPool2D(pool_size=(2, 2)))

# Input dimensions: (None, 48, 48, 32)
model.add(Convolution2D(64, (3,3), padding='same', use_bias=False))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())
# Input dimensions: (None, 48, 48, 64)
model.add(Convolution2D(64, (3,3), padding='same', use_bias=False))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())
model.add(MaxPool2D(pool_size=(2, 2)))

# Input dimensions: (None, 24, 24, 64)
model.add(Convolution2D(96, (3,3), padding='same', use_bias=False))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())
# Input dimensions: (None, 24, 24, 96)
model.add(Convolution2D(96, (3,3), padding='same', use_bias=False))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())
model.add(MaxPool2D(pool_size=(2, 2)))

# Input dimensions: (None, 12, 12, 96)
model.add(Convolution2D(128, (3,3),padding='same', use_bias=False))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())
# Input dimensions: (None, 12, 12, 128)
model.add(Convolution2D(128, (3,3),padding='same', use_bias=False))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())
model.add(MaxPool2D(pool_size=(2, 2)))
```

```

# Input dimensions: (None, 6, 6, 128)
model.add(Convolution2D(256, (3,3),padding='same',use_bias=False))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())
# Input dimensions: (None, 6, 6, 256)
model.add(Convolution2D(256, (3,3),padding='same',use_bias=False))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())
model.add(MaxPool2D(pool_size=(2, 2)))

# Input dimensions: (None, 3, 3, 256)
model.add(Convolution2D(512, (3,3), padding='same', use_bias=False))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())
# Input dimensions: (None, 3, 3, 512)
model.add(Convolution2D(512, (3,3), padding='same', use_bias=False))
model.add(LeakyReLU(alpha = 0.1))
model.add(BatchNormalization())

# Input dimensions: (None, 3, 3, 512)
model.add(Flatten())
model.add(Dense(512,activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(30))
model.summary()

```

conv2d_5 (Conv2D)	(None, 24, 24, 96)	02744
leaky_re_lu_5 (LeakyReLU)	(None, 24, 24, 96)	0
batch_normalization_5 (Batch Normalization)	(None, 24, 24, 96)	384
max_pooling2d_2 (MaxPooling2D)	(None, 12, 12, 96)	0
conv2d_6 (Conv2D)	(None, 12, 12, 128)	110592
leaky_re_lu_6 (LeakyReLU)	(None, 12, 12, 128)	0
batch_normalization_6 (Batch Normalization)	(None, 12, 12, 128)	512
conv2d_7 (Conv2D)	(None, 12, 12, 128)	147456
leaky_re_lu_7 (LeakyReLU)	(None, 12, 12, 128)	0
batch_normalization_7 (Batch Normalization)	(None, 12, 12, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 128)	0
conv2d_8 (Conv2D)	(None, 6, 6, 256)	294912
leaky_re_lu_8 (LeakyReLU)	(None, 6, 6, 256)	0
batch_normalization_8 (Batch Normalization)	(None, 6, 6, 256)	1024

conv2d_9 (Conv2D)	(None, 6, 6, 256)	589824
leaky_re_lu_9 (LeakyReLU)	(None, 6, 6, 256)	0
batch_normalization_9 (Batch Normalization)	(None, 6, 6, 256)	1024
max_pooling2d_4 (MaxPooling2D)	(None, 3, 3, 256)	0
conv2d_10 (Conv2D)	(None, 3, 3, 512)	1179648
leaky_re_lu_10 (LeakyReLU)	(None, 3, 3, 512)	0
batch_normalization_10 (Batch Normalization)	(None, 3, 3, 512)	2048
conv2d_11 (Conv2D)	(None, 3, 3, 512)	2359296
leaky_re_lu_11 (LeakyReLU)	(None, 3, 3, 512)	0
batch_normalization_11 (Batch Normalization)	(None, 3, 3, 512)	2048
flatten (Flatten)	(None, 4608)	0
dense (Dense)	(None, 512)	2359808
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 30)	15390
=====		
Total params: 7,268,670		
Trainable params: 7,264,318		
Non-trainable params: 4,352		

```
# Compile the model
```

```
model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae', 'acc'])
```

```
# Train the model
```

```
history = model.fit(train_images, train_keypoints, epochs=NUM_EPOCHS, batch_size=BATCH_SIZE,
```

Epoch 1/80

11/390 [.....] - ETA: 31:30 - loss: 1316.6342 - mae: 29.2299 -

KeyboardInterrupt

Traceback (most recent call last)

summarize history for mean_absolute_error

try:

```
plt.plot(history.history['mae'])
plt.plot(history.history['val_mae'])
plt.title('Mean Absolute Error vs Epoch')
plt.ylabel('Mean Absolute Error')
plt.xlabel('Epochs')
plt.legend(['train', 'validation'], loc='upper right')
plt.show()
# summarize history for accuracy
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('Accuracy vs Epoch')
plt.ylabel('Accuracy')
plt.xlabel('Epochs')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Loss vs Epoch')
plt.ylabel('Loss')
plt.xlabel('Epochs')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```

except:

```
print("One of the metrics used for plotting graphs is missing! See 'model.compile()'s `me
```

```
One of the metrics used for plotting graphs is missing! See 'model.compile()'s `metrics`
```

Fit the model on full dataset

```
model.fit(train_images, train_keypoints, epochs=NUM_EPOCHS, batch_size=BATCH_SIZE)
```

Epoch 1/80

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-25-b696fd2240bb> in <module>()
----> 1 model.fit(train_images, train_keypoints, epochs=NUM_EPOCHS,
batch_size=BATCH_SIZE)
```


Predicting on Test Set

```
quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
test_preds = model.predict(test_images)
----> 60                                inputs, attrs, num_outputs)
```

Visualizing Test Predictions

```
fig = plt.figure(figsize=(20,16))
for i in range(20):
    axis = fig.add_subplot(4, 5, i+1, xticks=[], yticks=[])
    plot_sample(test_images[i], test_preds[i], axis, "")
plt.show()
```





Generating Submission File



```
feature_names = list(idlookup_data['FeatureName'])
image_ids = list(idlookup_data['ImageId']-1)
row_ids = list(idlookup_data['RowId'])

feature_list = []
for feature in feature_names:
    feature_list.append(feature_names.index(feature))

predictions = []
for x,y in zip(image_ids, feature_list):
    predictions.append(test_preds[x][y])

row_ids = pd.Series(row_ids, name = 'RowId')
locations = pd.Series(predictions, name = 'Location')
locations = locations.clip(0.0,96.0)
submission_result = pd.concat([row_ids,locations],axis = 1)
submission_result.to_csv('submission.csv',index = False)
```

✓

4s

completed at 3:04 AM

●

✕